

# Human Activity Recognition with Deep Learning

## DNN vs CNN

Jainel Liz Boban

### Table of Contents

Load data . . . . .	2
Visualize data . . . . .	2
Train-validation-test splitting . . . . .	3
Model 1: DNN (2-hidden layer) + Early stoppping . . . . .	4
Plot for performance on training and validation data . . . . .	5
Performance evaluation on training and validation data . . . . .	6
Model 2: DNN (3-hidden layer) + Dropout + L2 Regularization with tuning for dropout rate, $\lambda$ , learning rate $\eta$ and no. of nodes in a layer. . . . .	7
Plot of predictive performance top 5 Tuned Model on training and validation data . . . . .	8
Extract runs with validation accuracy>95% and print . . . . .	11
Performance evaluation on training and validation data with model using optimal hyperpa- rameters . . . . .	12
Model 3: CNN . . . . .	13
Plot predictive performance of model on training and validaiton data . . . . .	13
Performance evaluation on training and validation data . . . . .	15
Choosing best fit model . . . . .	15
Best Model . . . . .	16
Retrain best model using full training data (train+val) . . . . .	16
Test performance overall . . . . .	16
Class-wise performance . . . . .	17

## Load data

- The data consists of `x`, `y`, `x_test` and `y_test` objects.
- The dimension of `x` and `x_test` are given below.
- The aim is to predict one of 19 activities from 125x45 multivariate data.

```
library(keras) # for DNN and CNN
# Load RData file and list objects in it
load("data_activity_recognition.RData")

dim(x)          # full training set input size
```

```
[1] 8170 125 45
```

```
dim(x_test)     # test set input size
```

```
[1] 950 125 45
```

- Each observation in '`x`' and '`x_test`' is a signal segment represented as a 2D matrix of 125 sampling instants by 45 sensor features.
- There are 8170 samples in training and 950 samples in the test set.
- Therefore, `x` and `x_test` are 3D arrays and can be used for Convolutional Neural network (CNN) using 1D convolutional layer or reshape it as 4D to use 2D convolutional layer.
- However, we can reshape it to 2D for use in Deep Neural Networks(DNN).

## Visualize data

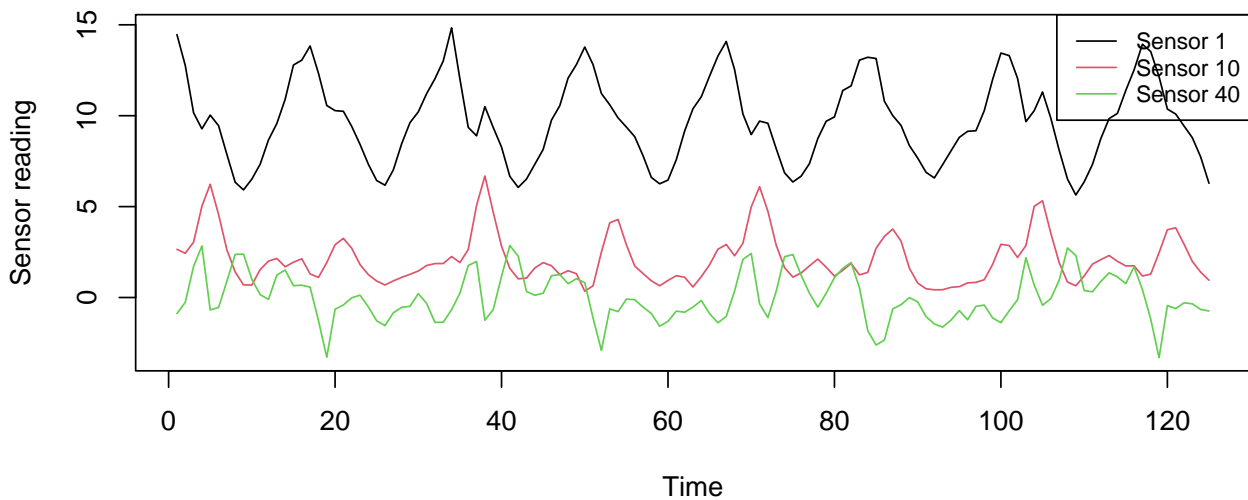
The sensor data can be viewed as a time series (sampling instants) as shown in plot:

```
par(mfrow = c(2, 1))
# custom plot function to view data
plot_signal <- function(index, data, sensors, title = "") {
  matplot(data[index,,sensors], type = "l", lty = 1,
          ylab = "Sensor reading", xlab = "Time",
          main = title, col = 1:length(sensors))
  legend("topright", legend = paste("Sensor", sensors),
        col = 1:length(sensors), lty = 1, cex = 0.8)
}

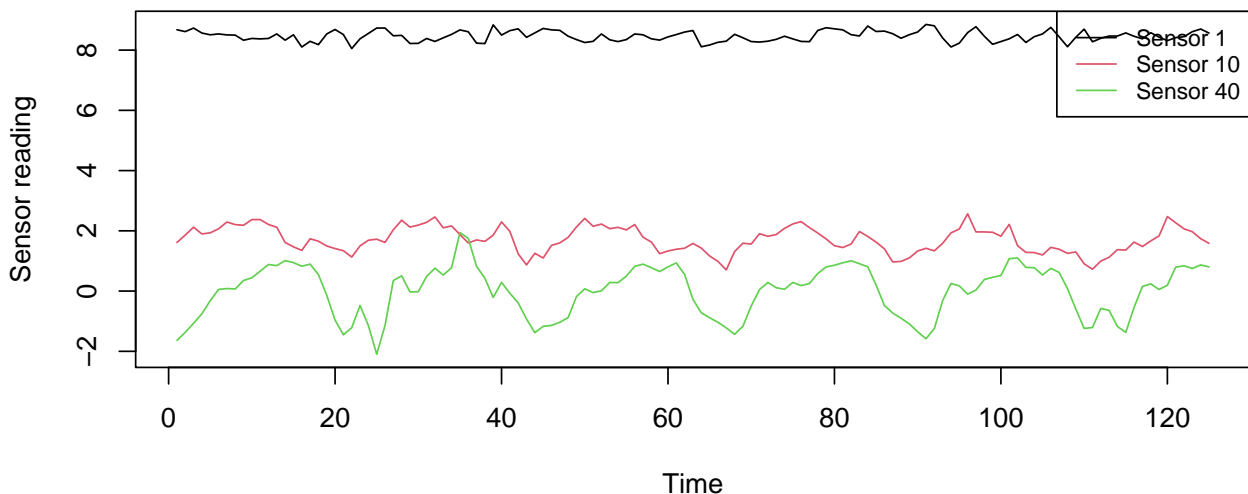
# first walking signal
index_walking <- which(grepl("walking", y))[1]
# print 1st, 10th and 40th sensor output for walking activity
plot_signal(index_walking, x, sensors = c(1, 10, 40),
            title = paste("Walking - Index", index_walking))

# first cycling signal
index_cycling <- which(grepl("cycling", y))[1]
plot_signal(index_cycling, x, sensors = c(1, 10, 40),
            title = paste("Cycling - Index", index_cycling))
```

## Walking – Index 2



## Cycling – Index 4



### Train-validation-test splitting

- We will split the entire training data into ‘train’ and ‘test’ sets as shown.
- We will pick same number of samples for validation data-set as in the test set (950).
- For DNNs, each input sample must be 1D feature vector. Therefore, the input data should be first converted to a 2D matrix format where each row corresponds to one observation and each column corresponds to one feature.
- Since we are doing multi-class classification, we convert target sets to one-hot encoding format using `to_categorical()` from `keras` library.

```
set.seed(24214925) # seed for reproducibility

# one-hot encoding of target variable
class_labels <- levels(factor(y))
n_classes <- length(class_labels) # number of classes to identify = 19
y <- as.numeric(factor(y)) - 1
y_test <- as.numeric(factor(y_test)) - 1
y <- keras::to_categorical(y, num_classes = n_classes)
y_test <- keras::to_categorical(y_test, num_classes = n_classes)
```

```
# convert train and test input into 125x45 vectors for DNN
x_dnn <- array_reshape(x, dim = c(nrow(x), 125*45))
x_test_dnn <- array_reshape(x_test, dim = c(nrow(x_test), 125*45))

#val train splitting
val <- sample(1:nrow(x), 950) # there are 950 activity labels in x_test
train <- setdiff(1:nrow(x), val) # remaining indices for training
y_val <- y[val,] # target val (DNN and CNN)
y_train <- y[-val,] # target train (DNN and CNN)
x_val_dnn <- x_dnn[val,] # input val (DNN)
x_train_dnn <- x_dnn[-val,] # input train (DNN)
x_val_cnn <- x[val,,] # input val (CNN)
x_train_cnn <- x[-val,,] # input train (CNN)

N = nrow(x) # no. of observation in training set (DNN and CNN)
V <- ncol(x_dnn) # (DNN)
```

### Model 1: DNN (2-hidden layer) + Early stopping

- The first model is a deep neural network with multiple layers- it has 2 hidden layers, both using “RELU”(Rectified Linear Unit) activation function.
- This activation function is chosen for the hidden layers because it avoids the vanishing gradient issue more effectively, as its derivative remains 1 for positive inputs unlike sigmoid or tanh.
- The output activation is selected as “Softmax” because the model is predicting one of many classes (19 activities) and it turns the outputs into probabilities for each class.
- This is well suited for the loss function “categorical cross-entropy” that we use here for multi-class classification.
- We use the “rmsprop” optimizer to use adaptive learning rate per parameter to adjust individual weights.
- It is chosen here as it is faster than Stochastic Gradient Descent (SGD) but simpler than Adam.
- Here, we use **early stopping** as a regularization technique to prevent over-fitting.
- This lets us to monitor model performance on the validation set during training (say “validation loss” or “validation accuracy”) and stop training when performance stops improving.
- *Patience* attribute allows training to continue for a number of epochs even if no improvement is seen.
- This gives the model a chance to recover in that period.
- According to below code the training will stop early if it does not improve for 20 continuous epochs, saving training time.
- Additionally, we also reduce learning rate by 10% if loss does not improve for 10 continuous epochs, giving the optimizer a smaller step size to tune better (fine-tune).
- This is external to RMSprop as it adjusts the global learning rate and works well together with it.

```
# Model 1 configuration
model1 <- keras_model_sequential() %>%
#first hidden layer with 128 hidden units and RELU activation
layer_dense(units = 128, activation = "relu", input_shape = V) %>%
#second hidden layer with 64 hidden units and RELU activation
layer_dense(units = 64, activation = "relu") %>%
#output layer with softmax activation with multiclass
layer_dense(units = ncol(y_train), activation = "softmax") %>%
compile( # compile model
  loss = "categorical_crossentropy", # loss function for multiclass
  metrics = "accuracy", # metric to evaluate
  optimizer = optimizer_rmsprop(), # adapt optimization rmsprop
)
# training and evaluation
fit_model1<- model1 %>% fit(
  x = x_train_dnn, y = y_train,
  validation_data = list(x_val_dnn, y_val),
  epochs = 100, # model goes through entire training set 100 times
  batch_size = round(N * 0.01), # Batch size = 1% of training set
  verbose = 0,
  callbacks = list(
    # early stopping if validation accuracy doesn't improve for 20 epochs
```

```

    callback_early_stopping(monitor = "val_accuracy", patience = 20),
    # learning rate reduced by 0.1 if loss do not improve for 10 epochs
    callback_reduce_lr_on_plateau(monitor = "loss", patience=10, factor=0.1))
)
# we can also count parameters in a model as shown
count_params(model1)

```

[1] 729619

### Plot for performance on training and validation data

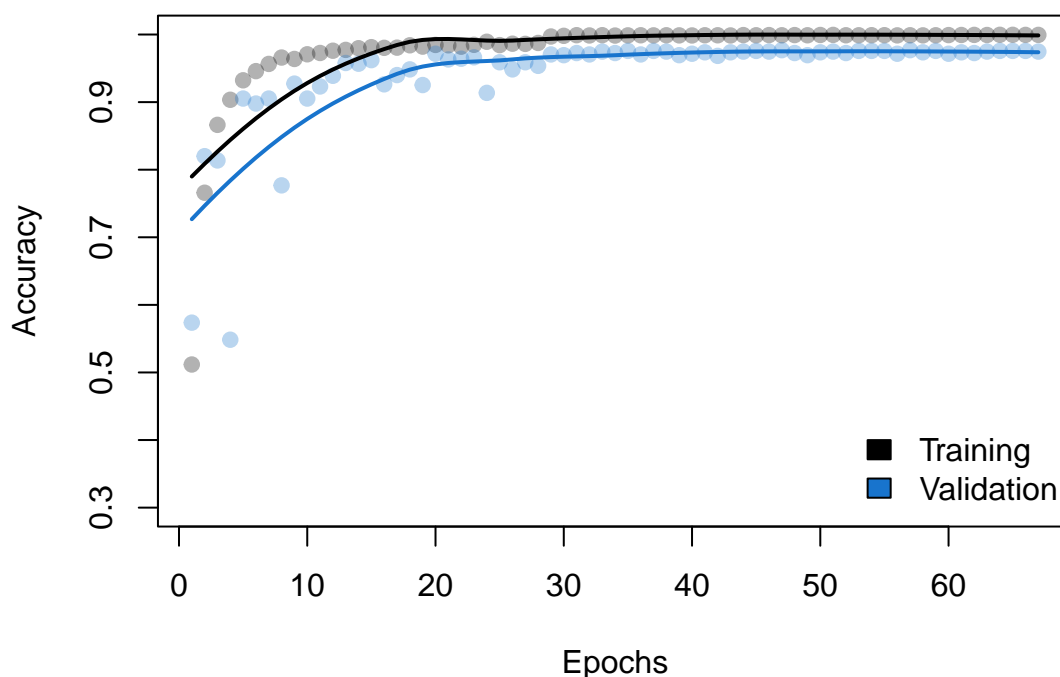
```

# function to add a smooth line to points in plot
smooth_line <- function(y , span=0.75) {
  x <- 1:length(y)
  out <- predict(loess(y ~ x), span=span)
  return(out)
}
cols <- c("black", "dodgerblue3")

# learning curves
out_model1 <- cbind(fit_model1$metrics$loss, fit_model1$metrics$val_loss,
                    fit_model1$metrics$accuracy, fit_model1$metrics$val_accuracy)

# accuracy plot
matplot(out_model1[,3:4], pch = 19, ylab = "Accuracy", xlab = "Epochs",
        col = adjustcolor(cols, 0.3), ylim = c(0.3, 1))
matlines(apply(out_model1[,3:4], 2, smooth_line), lty = 1, col = cols, lwd = 2)
legend("bottomright", legend = c("Training", "Validation"), fill = cols, bty = "n")

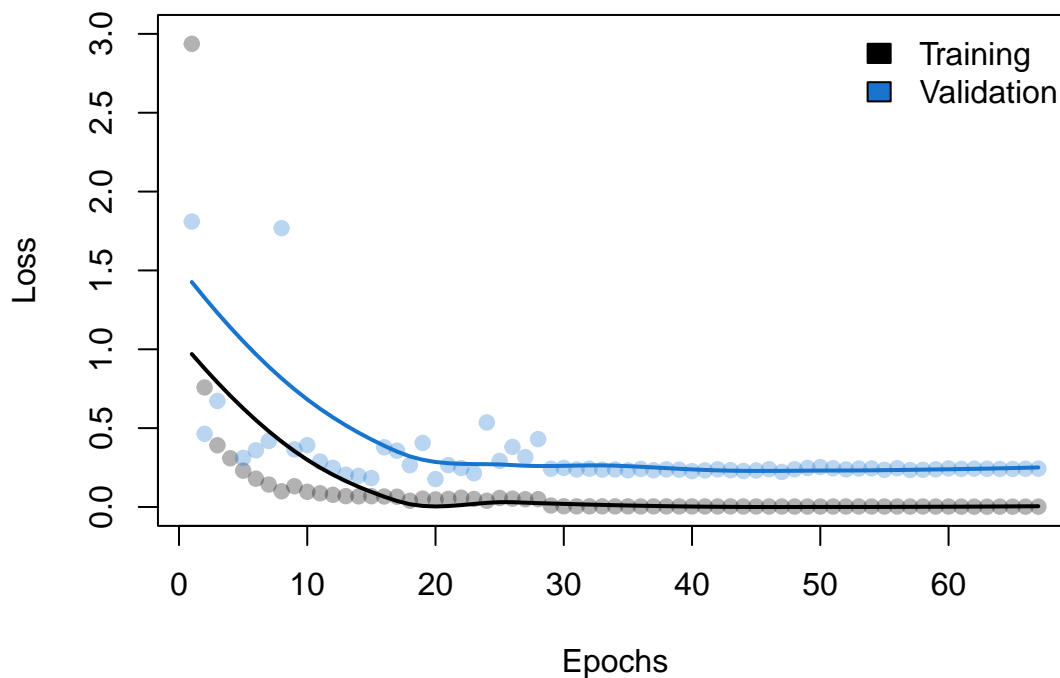
```



```

# loss plot
matplot(out_model1[,1:2], pch = 19, ylab = "Loss", xlab = "Epochs",
        col = adjustcolor(cols, 0.3), ylim = c(0, 3))
matlines(apply(out_model1[,1:2], 2, smooth_line), lty = 1, col = cols, lwd = 2)
legend("topright", legend = c("Training", "Validation"), fill = cols, bty = "n")

```



#### Observations:

- There seems to be good generalization with no significant over-fitting as the training and validation curves are close (training accuracy almost saturates close to 1).
- We can also see that the training stopped when a certain number of epochs were reached (did not reach 100 as specified) before significant over-fitting.
- There is a gap between training and validation loss which shows a slight over-fitting. However, it is not that severe and we can say that the model has learned the patterns without just simply learning the training data.

#### Performance evaluation on training and validation data

- We can assess the performance of the model on the training and validation sets at convergence using `evaluate()` as shown:

```
#performance on training data
train_metrics_model1 <- model1 %>% evaluate(x_train_dnn, y_train, verbose = 0)
train_metrics_model1
```

```
      loss      accuracy
0.001832718 0.999445975
```

```
# evaluate performance on validation data
val_metrics_model1 <- model1 %>% evaluate(x_val_dnn, y_val, verbose = 0)
val_metrics_model1
```

```
      loss      accuracy
0.2430505 0.9747369
```

#### Observations:

- The final training accuracy is 99.94% with a very low training loss of 0.0018.
- Therefore the model fits the training data really well.
- The validation accuracy is 97.47% which is also high.
- The slightly high validation loss of 0.2431 suggests a small amount of over-fitting, but it is not as severe.
- Overall, the model generalizes well on unseen data (validation data) and the early stopping and regularization method seems to be effective.

**Model 2: DNN (3-hidden layer) + Dropout + L2 Regularization with tuning for dropout rate,  $\lambda$ , learning rate  $\eta$  and no. of nodes in a layer.**

- Below is a training code stored in an external R file. This script will be passed to *tuning\_run()* using *tfruns* as shown in the next code.
- *tuning\_run()* re-executes the model code for each combination of hyperparameters.
- In this deep neural network model, we will use 3 hidden layers with “RELU” activation function (same activation function as before for hidden layers).
- Output activation function is selected again as “Softmax”.
- We use FLAGS here to define tunable hyperparameters such as: dropout rate, regularization parameter (L2 regularization), learning rate, batch size and first hidden layer size.
- We use the “Adam” optimizer here as it is a fast and widely used adaptive learning rate method.
- **Dropout** technique use used here for regularization to prevent overfitting.
- This will randomly drop some units during training time, so that the model will not rely on specific units (neurons).
- As dropout rate increases, there will be more regularization. However, we have to also prevent under-fitting. Hence, we can tune this parameter to get good performance.
- In addition to dropout, we can also use **L2 regularization** which adds a penalty term to the loss function to reduce complexity as shown:

$$E_Q(\mathbf{w}, \lambda) = E(\mathbf{w}) + \lambda \frac{1}{2} \mathbf{w}^\top \mathbf{w}$$

- Tuning penalty parameter  $\lambda$  will allow us to balance between under and over-fitting.
- Weight update equation with L2 regularization becomes:

$$\mathbf{w} \leftarrow \mathbf{w}(1 - \eta\lambda) - \eta \nabla E(\mathbf{w})$$

- As we can see **Learning rate**  $\eta$  controls how fast the model learns and also affects the optimizer.
- Therefore, by well-tuning  $\eta$ , we can get faster convergence.
- Tuning helps to balance between overshooting (too high  $\eta$ ) and slow learning (too low  $\eta$ ).
- The **batch-size** is also made tunable in this model.
- There should be a balance for batch size as too small batch size could lead to slow training and too big batch size may overfit the training data.
- The **number of neurons/units** in the layers of the network can also be tuned as it decides the complexity of the functions it can learn.
- Too many neurons may lead to overfitting and too few may underfit.
- Due to computational limitation, this has been demonstrated only for hidden layer 1 in the code:

```
# content of file 'Model2_tfruns.R'

# default flags for the code
FLAGS <- flags(
  flag_numeric("dropout", 0.3),
  flag_numeric("lambda", 0.01),
  flag_numeric("lr", 0.01),
  flag_numeric("bs", 100),
  flag_integer("size_1", 256)
)

# model configuration
model_tune <- keras_model_sequential() %>%
  # first hidden layer units tunable
  layer_dense(units = FLAGS$size_1, input_shape = V, activation = "relu",
    name = "layer_1",
    # L2 Regularization done with tunable parameter lambda
    kernel_regularizer = regularizer_l2(FLAGS$lambda)) %>%
  # dropout rate tunable
  layer_dropout(rate = FLAGS$dropout) %>%
  layer_dense(units = 128, activation = "relu", name = "layer_2",
    kernel_regularizer = regularizer_l2(FLAGS$lambda)) %>%
  layer_dropout(rate = FLAGS$dropout) %>%
  layer_dense(units = 64, activation = "relu", name = "layer_3",
    kernel_regularizer = regularizer_l2(FLAGS$lambda)) %>%
```

```

layer_dropout(rate = FLAGS$dropout) %>%
layer_dense(units = ncol(y_train), activation = "softmax",
             name = "layer_out") %>%
compile(loss = "categorical_crossentropy", metrics = "accuracy",
        # adaptive optimizer Adam with tunable learning rate
        optimizer = optimizer_adam(learning_rate = FLAGS$lr),
)
# training and evaluation
fit_model_tune <- model_tune %>% fit(
  x = x_train_dnn, y = y_train,
  validation_data = list(x_val_dnn, y_val),
  epochs = 100,
  batch_size = FLAGS$bs, # batch size tunable
  verbose = 0,
  callbacks = callback_early_stopping(monitor = "val_accuracy",
                                     patience = 20))

```

- Using the above file (saved as *Model2\_tfruns.R* in the current working directory) can be used along package *tfruns* to easily tune and experiment with different settings (hyperparameter combinations) for the model.
- We specify a grid of values for the hyperparameters of interest.
- We will be using random search because in case of grid search, it uses all combinations (3x3x3x2x2=108 combinations in total as per below grid) for tuning and is computationally expensive.
- It is also better than manually tuning hyperparameters as it takes more time to manually tune a couple of parameters and may miss best combinations by only testing fewer options.
- In random search, we randomly sample a proportion of the total combinations (20% as per below code, reducing the count to much lesser 22 combinations).
- This will help us find acceptable results with much fewer runs.

```

set.seed(24214925)

library(tfruns)
# grid with different combinations for tuning
dropout_set <- c(0, 0.3, 0.4) # dropout rate
lambda_set <- c(0, 0.005, 0.01) # regularization parameter
lr_set <- c(0.001, 0.005, 0.01) # learning rate
bs_set <- c(0.01, 0.05)*N # batch size in percentages of no. of samples
size1_set <- c(512, 256) # tuning no. of nodes first hidden layer

# Use random search to pick random combinations from 108 combinations above
invisible(runs <- tuning_run("Model2_tfruns.R",
runs_dir = "runs_model2", # save results in folder
flags = list(
  dropout = dropout_set,
  lambda = lambda_set,
  lr = lr_set,
  bs = bs_set,
  size_1 = size1_set),
sample = 0.2)) # sample 20% = 22 combinations

```

## Plot of predictive performance top 5 Tuned Model on training and validation data

```

#custom function to read training matrix
read_metrics <- function(path, files = NULL){
  path <- paste0(path, "/")
  if (is.null(files)) files <- list.files(path)
  n <- length(files) # no. of runs (folders)
  out <- vector("list", n) # store metrics for each run
  for (i in 1:n){ # loop over each folder
    dir <- paste0(path, files[i], "/tfruns.d/") # path to directory
    tryCatch({
      #training metrics
      out[[i]] <- jsonlite::fromJSON(paste0(dir, "metrics.json"))
    }, error = function(e){
      #validation metrics
    })
  }
}

```



```

# hyperparameter values
out[[i]]$flags <- jsonlite::fromJSON(paste0(dir, "flags.json"))
#evaluation metrics
out[[i]]$evaluation <- jsonlite::fromJSON(paste0(dir, "evaluation.json"))
}, error = function(e){
  # skip corrupt files
  message("Corrupt file ", files[i], ": ", e$message)
  out[[i]] <- NULL
})
}
out <- Filter(Negate(is.null),out) # remove failed runs
return(out)
}
# results from folders
out <- read_metrics("runs_model2")
# training and validation across runs
train_acc <- sapply(out, "[[", "accuracy")
val_acc <- sapply(out, "[[", "val_accuracy")
train_loss <- sapply(out, "[[", "loss")
val_loss <- sapply(out, "[[", "val_loss")

```

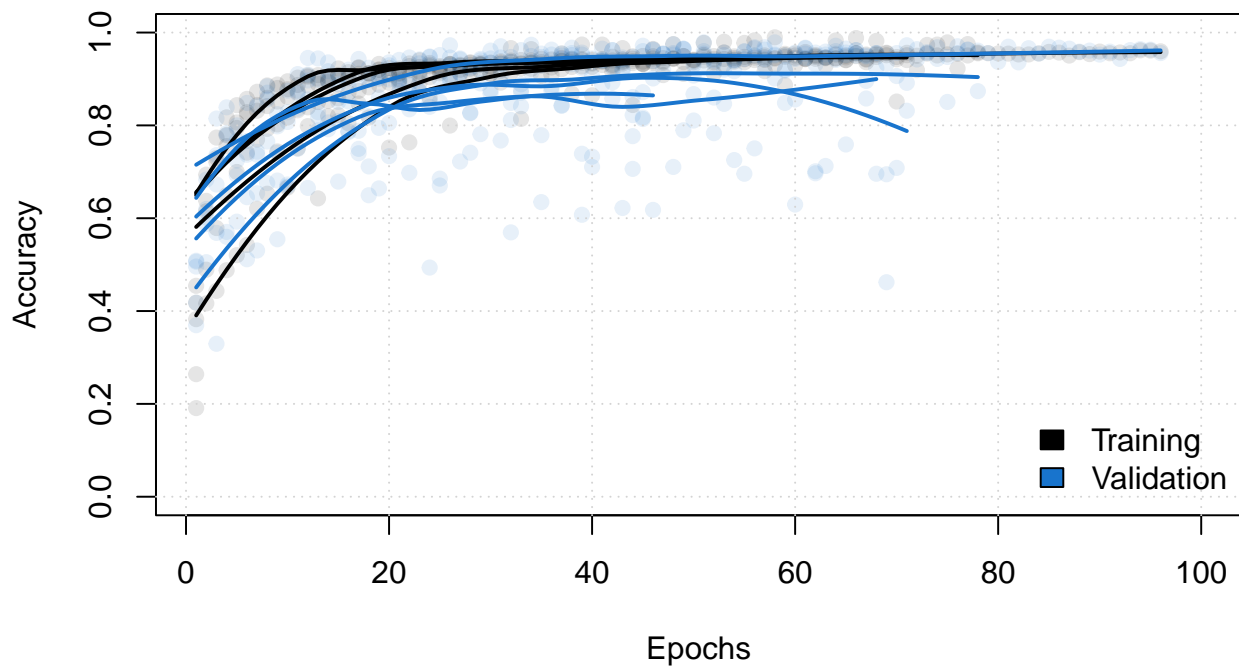
Accuracy and loss curves are given below:

```

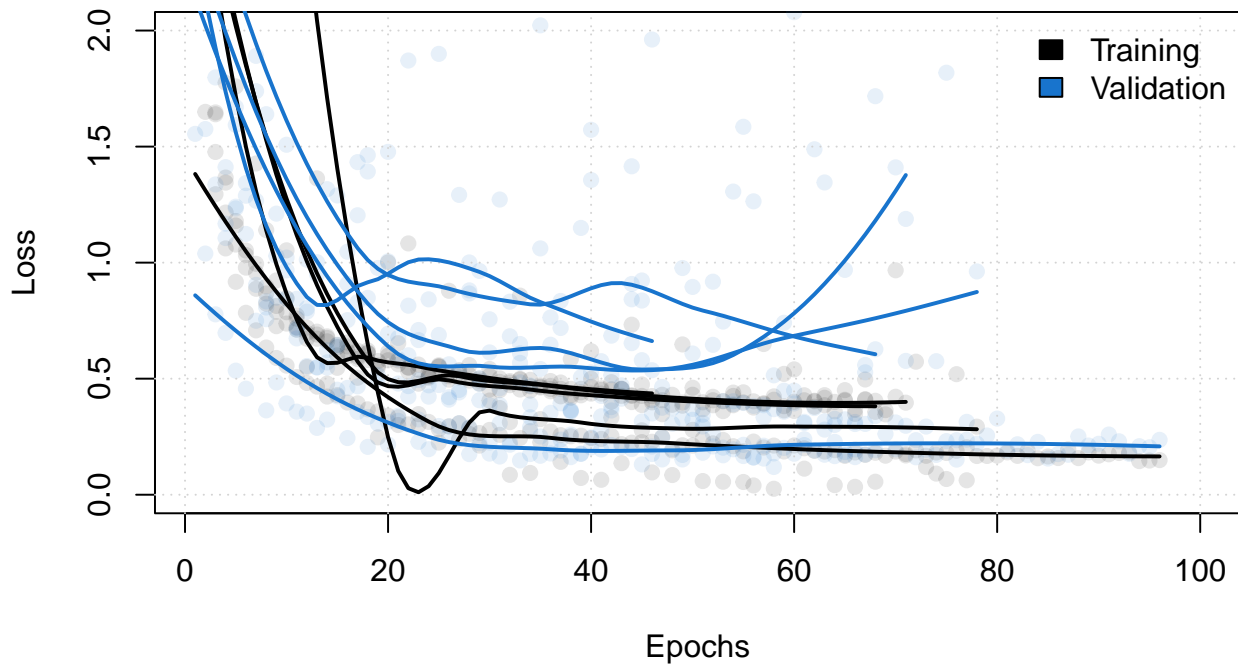
# select top 5 runs by validation accuracy for plotting
sel <- 5
top <- order(apply(val_acc, 2, max, na.rm = TRUE), decreasing = TRUE)[1:sel]
cols_m2 <- rep(c("black", "dodgerblue3"), each = sel)

# plot accuracy curves
out_acc <- cbind(train_acc[, top], val_acc[, top])
matplot(out_acc, pch = 19, ylab = "Accuracy", xlab = "Epochs",
  col = adjustcolor(cols_m2, 0.1), ylim = c(0, 1))
grid()
tmp <- apply(out_acc, 2, smooth_line, span=0.8) # smoothed lines over curves
tmp <- sapply(tmp, "length<=", 100) # 100 epochs by default
matlines(tmp, lty = 1, col = cols_m2, lwd = 2)
legend("bottomright", legend = c("Training", "Validation"), # identify train/test
  fill = unique(cols_m2), bty = "n")

```



```
# plot loss curves
out_loss <- cbind(train_loss[, top], val_loss[, top])
matplot(out_loss, pch = 19, ylab = "Loss", xlab = "Epochs",
        col = adjustcolor(cols_m2, 0.1), ylim = c(0, 2))
grid()
tmp <- apply(out_loss, 2, smooth_line, span = 0.8)
tmp <- sapply(tmp, "length<-", 100)
matlines(tmp, lty = 1, col = cols_m2, lwd = 2)
legend("topright", legend = c("Training", "Validation"),
      fill = unique(cols_m2), bty = "n")
```



#### Observations:

- The training accuracy is high while validation accuracy is lower and has more variability.
- This could be due to overfitting in some runs.
- However, we can also see runs that reach high validation accuracy above 90% showing that tuning helped find effective hyperparameter combinations from the available sample.
- It can be also observed that the loss curves are highly fluctuating (even after smoothing).
- The gap between training and validation loss curves suggests the model is more uncertain when predicting on unseen data.
- From the plots, Model 1 seems like a better fit compared to Model 2.

#### Extract runs with validation accuracy > 95% and print

The below code can be used to view the top runs (hyperparameter combinations) that achieved the maximum validation accuracy during training.

```
res <- ls_runs(metric_val_accuracy > 0.95, runs_dir = "runs_model2",
              order = metric_val_accuracy)
colu <- c("metric_val_accuracy", grep("flag", colnames(res), value = TRUE),
          "epochs_completed")
res[1:5, colu]
```

Data frame: 5 x 7

	metric_val_accuracy	flag_dropout	flag_lambda	flag_lr	flag_bs	flag_size_1
1	0.9726	0.0	0.005	0.001	81.7	512
2	0.9558	0.3	0.000	0.001	81.7	256
3	0.9547	0.0	0.010	0.001	408.5	512
4	0.9526	0.0	0.010	0.001	408.5	512
NA	NA	NA	NA	NA	NA	NA

	epochs_completed
1	71
2	96
3	100
4	100
NA	NA

## Performance evaluation on training and validation data with model using optimal hyperparameters

- The first row of 'res' has the highest metric\_val\_accuracy of %.
- It also ran for the full 100 epochs, implying that early stopping was not triggered.
- This suggests that the model continued improving, hence we will use the same parameter values in this first combination to fit Model 2 using a slightly higher number of epochs.
- It can be noticed that in this best combination, the dropout rate is 0.

```
# final deployment using optimal hyperparameters at res[1,]
model2_optimal <- keras_model_sequential() %>%
  layer_dense(units = res$flag_size_1[1], input_shape = V, activation = "relu",
              name = "layer_1",
              kernel_regularizer = regularizer_l2(res$flag_lambda[1])) %>%
  layer_dropout(rate = res$flag_dropout[1]) %>%
  layer_dense(units = 128, activation = "relu", name = "layer_2",
              kernel_regularizer = regularizer_l2(res$flag_lambda[1])) %>%
  layer_dropout(rate = res$flag_dropout[1]) %>%
  layer_dense(units = 64, activation = "relu", name = "layer_3",
              kernel_regularizer = regularizer_l2(res$flag_lambda[1])) %>%
  layer_dropout(rate = res$flag_dropout[1]) %>%
  layer_dense(units = ncol(y_train), activation = "softmax", name = "layer_out") %>%
  compile(loss = "categorical_crossentropy", metrics = "accuracy",
          optimizer = optimizer_rmsprop(learning_rate = res$flag_lr[1]),
  )
# training and evaluation
fit_model2_optimal <- model2_optimal %>% fit(
  x = x_train_dnn, y = y_train,
  validation_data = list(x_val_dnn, y_val),
  epochs = 150,
  batch_size = res$flag_bs[1],
  verbose = 0,
  callbacks = callback_early_stopping(monitor = "val_accuracy",
                                     patience = 20))
#count parameters
count_params(model2_optimal)
```

```
[1] 2955667
```

```
#performance on training data
train_metrics_model2 <- model2_optimal %>% evaluate(x_train_dnn, y_train, verbose = 0)
train_metrics_model2
```

```
      loss accuracy
0.3274313 0.9635734
```

```
# evaluate performance on validation data
val_metrics_model2 <- model2_optimal %>% evaluate(x_val_dnn, y_val, verbose = 0)
val_metrics_model2
```

```
      loss accuracy
0.4531380 0.9284211
```

### Observations:

- The training accuracy reached 96.36%.
- However, the validation accuracy is 92.84% which is lower.
- Validation loss is around 0.4531 compared to training loss of 0.3274.
- There seems to be some over-fitting as the model works well on training but not as much (or close) for validation.
- The total parameter count is quite high (we used more layers) which could contribute to this overfitting. - Overall, the performance is strong, but there is still some room for improvement with better tuning parameters or using other regularization techniques.

### Model 3: CNN

- We use a 1D Convolutional Neural Network (CNN) since our input is a 3D tensor, where each sample consists of a 2D feature matrix over time (we do not have other spatial dimensions like an image).
- Hence the kernel size and pooling size will also be 1D.
- The network consist of 2 1D-convolution layers interleaved by 2 1D-max-pooling layers.
- A kernel size of 5 (say like first convolutional layer) means each filter views 5 features at a time (default stride is 1).
- A max pooling with window size 2 has a default stride of 2.
- A convolutional layer helps learn local signal curve features using small filters.
- Pooling prevents overfitting by by reducing spatial size keeping important information.
- Flattening before dense layers converts the final 3D tensor output into a 1D vector to be def into fully connected layers.
- Dropout regularizations (30-40%) have been applied between convolutional layers as well as fully connected layers to prevent overfitting and improve generalization.
- Adam optimizer has been used for this model.

```
# model 3 configuration
model3 <- keras_model_sequential() %>%
  # 1D convolutional layer with 5*1 kernel_size, strides=1 by default
  layer_conv_1d(filters = 128, kernel_size = 5, activation = "relu",
    input_shape = c(125, 45)) %>% # input shape 3D
  # max pooling strides = pool_size = 2
  layer_max_pooling_1d(pool_size = 2) %>%
  # 40% dropout after this layer
  layer_dropout(0.4) %>%
  layer_conv_1d(filters = 128, kernel_size = 3, activation = "relu") %>%
  layer_max_pooling_1d(pool_size = 2) %>%
  layer_dropout(0.3) %>%
  layer_flatten() %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dropout(0.3) %>%
  layer_dense(units = n_classes, activation = "softmax") %>%
  compile(loss = "categorical_crossentropy", metrics = "accuracy",
    optimizer= optimizer_adam()) # adam optimizer

fit_model3 <- model3 %>% fit(
  x = x_train_cnn, y = y_train,
  validation_data = list(x_val_cnn, y_val),
  epochs = 100,
  batch_size = 72,
  verbose = 0
)
# performance on training data
model3 %>% evaluate(x_train_cnn, y_train, verbose = 0)
```

```
      loss    accuracy
0.01156339 0.99626040
```

```
# evaluate performance on validation data
model3 %>% evaluate(x_val_cnn, y_val, verbose = 0)
```

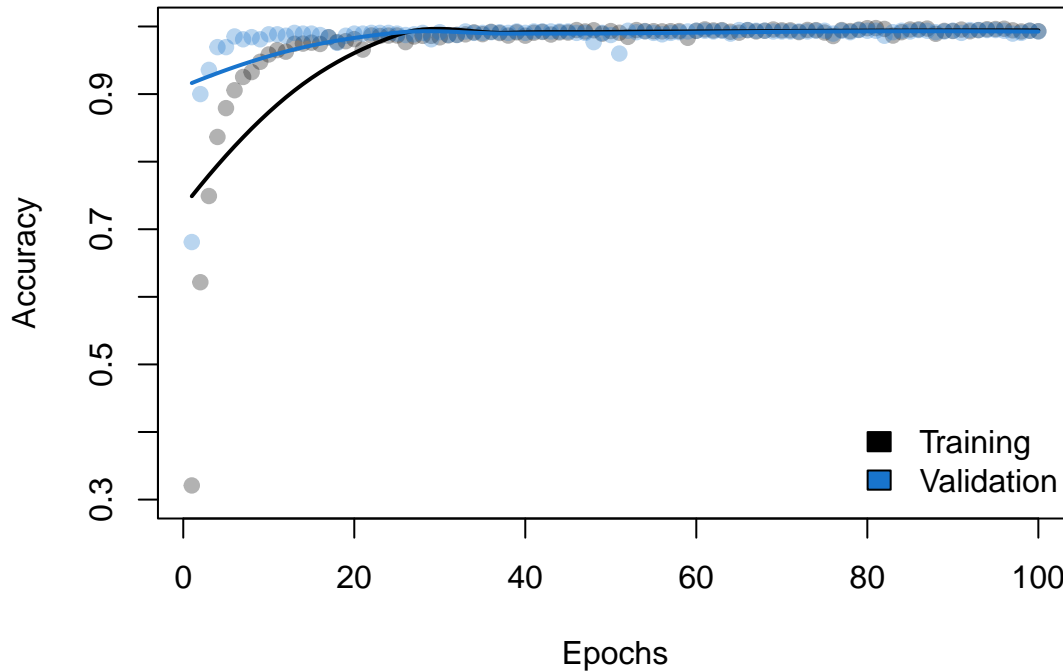
```
      loss    accuracy
0.03033537 0.99263155
```

### Plot predictive performance of model on training and validation data

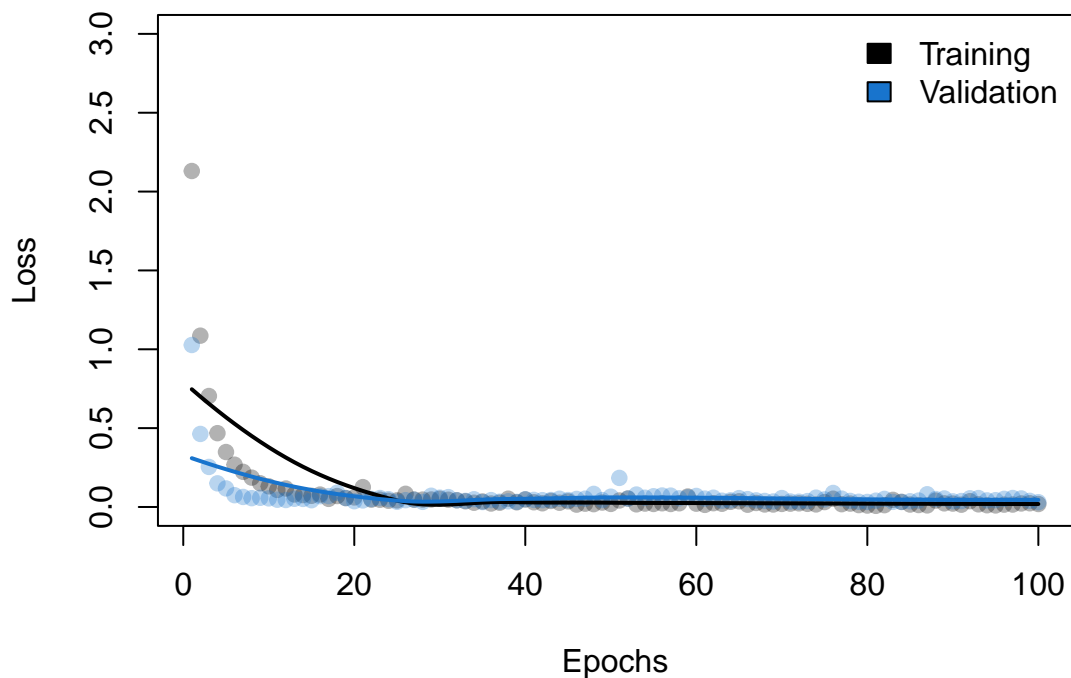
```
# learning curves
out_model3 <- cbind(fit_model3$metrics$loss, fit_model3$metrics$val_loss,
  fit_model3$metrics$accuracy, fit_model3$metrics$val_accuracy)

# accuracy plot
matplot(out_model3[,3:4], pch = 19, ylab = "Accuracy", xlab = "Epochs",
  col = adjustcolor(cols, 0.3), ylim = c(0.3, 1))
```

```
matlines(apply(out_model3[,3:4], 2, smooth_line), lty = 1, col = cols, lwd = 2)
legend("bottomright", legend = c("Training", "Validation"), fill = cols, bty = "n")
```



```
# loss plot
matplot(out_model3[,1:2], pch= 19, ylab= "Loss", xlab= "Epochs",
        col= adjustcolor(cols, 0.3), ylim= c(0, 3))
matlines(apply(out_model3[,1:2], 2, smooth_line), lty = 1, col = cols, lwd = 2)
legend("topright", legend = c("Training", "Validation"), fill = cols, bty = "n")
```



#### Observations:

- The training and validation curves almost overlap (for both training and validation).
- This shows that the model has very good generalization and has no overfitting.
- The pooling layers and dropout seem to have helped maintain generalization.
- It can be also seen that both accuracy reached a stable state obtaining high values quickly than other models.
- Both loss curves converge close to zero around 30-40 epochs and has no major fluctuations like before (model is consistent).
- Overall, the CNN model performs well both in terms of accuracy and loss (much better than the DNN models 1 & 2).

#### Performance evaluation on training and validation data

```
#performance on training data
train_metrics_model3 <- model3 %>% evaluate(x_train_cnn, y_train, verbose = 0)
train_metrics_model3
```

```
      loss  accuracy
0.01156339 0.99626040
```

```
# evaluate performance on validation data
val_metrics_model3 <- model3 %>% evaluate(x_val_cnn, y_val, verbose = 0)
val_metrics_model3
```

```
      loss  accuracy
0.03033537 0.99263155
```

#### Observations:

- Just as we had seen from plot, we can see the values of accuracies for both training and validation sets are high for model 3 using CNN (Training accuracy of 99.63% and validation accuracy of 99.26%)
- The losses are also low (training loss of 0.0116 and validation loss of 0.0303)
- Therefore, model 3 generalizes well to unseen data (no overfitting).

#### Choosing best fit model

Metric	Model 1	Model 2	Model 3
Training Accuracy	99.94%	96.36%	99.63%
Validation Accuracy	97.47%	92.84%	99.26%
Training Loss	0.0018	0.3274	0.0116
Validation Loss	0.2431	0.4531	0.0303

#### Observations:

- As discussed before, clearly, **Model 3 using CNN is the best model** among the 3.
- Model 3 balances both training and validation metrics.
- Although training loss is low for Model 1, we can see that Model 3 (CNN) gives the best predictive performance on validation data in terms of both accuracy and loss.
- This could be due to CNN's ability to preserve spatial structure and learn local patterns more effectively than a DNN.
- Max-pooling helps to keep the important features and also reduce complexity.

### Best Model

The best model found based on performance on validation data is **Model 3 (CNN)**.

#### Retrain best model using full training data (train+val)

```
model3_final <- keras_model_sequential() %>%
  layer_conv_1d(filters = 128, kernel_size = 5, activation = "relu",
               input_shape = c(125, 45)) %>%
  layer_max_pooling_1d(pool_size = 2) %>%
  layer_dropout(0.4) %>%
  layer_conv_1d(filters = 128, kernel_size = 3, activation = "relu") %>%
  layer_max_pooling_1d(pool_size = 2) %>%
  layer_dropout(0.3) %>%
  layer_flatten() %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dropout(0.3) %>%
  layer_dense(units = n_classes, activation = "softmax") %>%
  compile(loss = "categorical_crossentropy", metrics = "accuracy",
          optimizer= optimizer_adam()) # adam optimizer

fit_model3_final <- model3_final %>% fit(
  x = x, y = y,
  epochs = 100,
  batch_size = 72,
  verbose=0
)
```

#### Test performance overall

```
set.seed(24214925)
# overall test accuracy and loss
final_train_metrics <- model3_final %>% evaluate(x_test, y_test, verbose= 0)
final_train_metrics

      loss      accuracy
0.03109396 0.99368423
```

#### Observation:

- Very good test performance with 99.37% test accuracy and 0.0311 overall test loss.
- CNN model has good ability to distinguish between the 19 physical activities on unseen test data.
- Hence, it can generalize well.



## Class-wise performance

```
# find class with the highest predicted probability
test_hat <-model3_final %>% predict(x_test, verbose=0)%>% max.col()
tab <-table(max.col(y_test),test_hat) # confusion matrix

# class-accuracy/sensitivity = TP / Total in class
cl.acc <- diag(tab) / rowSums(tab)
# Precision = TP / (TP + FP)
precision <- diag(tab) / colSums(tab)

# print confusion matrix with results
cbind(tab, Acc=round(cl.acc,2), Prec=round(precision,2))
```

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	Acc	Prec
1	50	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1.00	1.00
2	0	50	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1.00	1.00
3	0	0	50	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1.00	1.00
4	0	0	0	50	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1.00	1.00
5	0	0	0	0	50	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1.00	1.00
6	0	0	0	0	0	50	0	0	0	0	0	0	0	0	0	0	0	0	0	1.00	0.96
7	0	0	0	0	0	0	50	0	0	0	0	0	0	0	0	0	0	0	0	1.00	1.00
8	0	0	0	0	0	0	0	50	0	0	0	0	0	0	0	0	0	0	0	1.00	1.00
9	0	0	0	0	0	0	0	0	50	0	0	0	0	0	0	0	0	0	0	1.00	1.00
10	0	0	0	0	0	2	0	0	0	44	0	0	1	3	0	0	0	0	0	0.88	1.00
11	0	0	0	0	0	0	0	0	0	0	50	0	0	0	0	0	0	0	0	1.00	1.00
12	0	0	0	0	0	0	0	0	0	0	0	50	0	0	0	0	0	0	0	1.00	1.00
13	0	0	0	0	0	0	0	0	0	0	0	0	50	0	0	0	0	0	0	1.00	0.98
14	0	0	0	0	0	0	0	0	0	0	0	0	0	50	0	0	0	0	0	1.00	0.94
15	0	0	0	0	0	0	0	0	0	0	0	0	0	0	50	0	0	0	0	1.00	1.00
16	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	50	0	0	0	1.00	1.00
17	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	50	0	0	1.00	1.00
18	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	50	0	1.00	1.00
19	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	50	1.00	1.00

```
# print results using labels
final_class_metrics <- data.frame(Class_No = 1:length(class_labels),Class_Label = class_labels,
  Accuracy = round(cl.acc,2),Precision = round(precision,2))
print(final_class_metrics, row.names = FALSE)
```

Class_No	Class_Label	Accuracy	Precision
1	asc_stairs	1.00	1.00
2	basketball	1.00	1.00
3	cross_trainer	1.00	1.00
4	cycling_horiz	1.00	1.00
5	cycling_vert	1.00	1.00
6	desc_stairs	1.00	0.96
7	jumping	1.00	1.00
8	lying_back	1.00	1.00
9	lying_side	1.00	1.00
10	moving_elevator	0.88	1.00
11	rowing	1.00	1.00
12	running_treadmill	1.00	1.00
13	sitting	1.00	0.98
14	stand_elevator	1.00	0.94
15	standing	1.00	1.00
16	stepper	1.00	1.00
17	walking	1.00	1.00
18	walking_tread	1.00	1.00
19	walking_tread_incl	1.00	1.00

## Observations:

- The final model using CNN shows excellent classification with 100% accuracy and precision for most activities.

- The high class-wise accuracy (sensitivity) shows that the model has learned to distinguish between the majority of activities effectively and is able to generalize well on unseen data.
- The high precision for most classes shows that when model predicted a certain activity it was usually correct.
- There are only 4 classes where either accuracy or precision is less than 1 with very few misclassifications as shown in confusion matrix.